

```

# =====
# FULL PIX2PIX LAB IMPLEMENTATION
# =====
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader, Dataset
import matplotlib.pyplot as plt
import numpy as np
import cv2

# =====
# CONFIG
# =====
BATCH_SIZE = 64
EPOCHS = 20
LR = 0.0002
LAMBDA_L1 = 100
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"
DEVICE = "cuda" if torch.cuda.is_available() else "cpu"

print("Running on: " + DEVICE)

# =====
# DATASET (Simulated Edges2Shoes)
# =====
class SimulatedEdges2ShoesDataset(Dataset):
    def __init__(self, root='./data', train=True):
        self.root = root
        self.train = train
        self.transform = transforms.Compose([
            transforms.Resize(256),
            transforms.CenterCrop(256),
            transforms.ToTensor(),
            transforms.Normalize((0.5,), (0.5,))])
    def __len__(self):
        return len(self.data)
    def __getitem__(self, idx):
        real_img_pil = self.data[idx]
        real_img = self.transform(real_img_pil)
        img_np = np.array(real_img_pil)
        edges = cv2.Canny(img_np, 100, 200)
        edges = transforms.ToPILImage()(edges)
        edge_img = self.transform(edges)
        return edge_img, real_img

train_ds = SimulatedEdges2ShoesDataset(train=True)
train_loader = DataLoader(train_ds, batch_size=BATCH_SIZE, shuffle=True)

# =====
# U-NET GENERATOR
# =====
class UNetGenerator(nn.Module):
    def __init__(self):
        super().__init__()
        self.enc1 = self.conv_block(1, 64, bn=False)
        self.enc2 = self.conv_block(64, 128)
        self.enc3 = self.conv_block(128, 256)
        self.enc4 = self.conv_block(256, 512)
        self.bottleneck = self.conv_block(512, 512)
        self.dec4 = self.up_block(512, 512)
        self.dec3 = self.up_block(1024, 256)
        self.dec2 = self.up_block(512, 128)
        self.dec1 = self.up_block(256, 64)
        self.final = nn.Sequential(
            nn.ConvTranspose2d(128, 1, 4, 2, 1),
            nn.Tanh()
        )
    def conv_block(self, in_c, out_c, bn=True):
        layers = [nn.Conv2d(in_c, out_c, 4, 2, 1, bias=False)]
        if bn:
            layers.append(nn.BatchNorm2d(out_c))
        layers.append(nn.LeakyReLU(0.2))
        return nn.Sequential(*layers)
    def up_block(self, in_c, out_c):
        return nn.Sequential(
            nn.ConvTranspose2d(in_c, out_c, 4, 2, 1, bias=False),
            nn.BatchNorm2d(out_c),
            nn.ReLU()
        )
    def forward(self, x):
        e1 = self.enc1(x)
        e2 = self.enc2(e1)
        e3 = self.enc3(e2)
        e4 = self.enc4(e3)
        b = self.bottleneck(e4)
        d4 = self.dec4(b)
        d3 = torch.cat((d4, e3), dim=1)
        d2 = self.dec3(d3)
        d2 = torch.cat((d2, e2), dim=1)
        d1 = self.dec1(d2)
        d1 = torch.cat((d1, e1), dim=1)
        return self.final(d1)

# =====
# PATCHGAN DISCRIMINATOR
# =====
class PatchGANDiscriminator(nn.Module):
    def __init__(self):
        super().__init__()
        self.model = nn.Sequential(
            nn.Conv2d(2, 64, 4, 2, 1),
            nn.LeakyReLU(0.2),
            nn.Conv2d(64, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128),
            nn.LeakyReLU(0.2),
            nn.Conv2d(128, 256, 4, 2, 1, bias=False),
            nn.BatchNorm2d(256),
            nn.LeakyReLU(0.2),
            nn.Conv2d(256, 1, 4, 1, 1)
        )
    def forward(self, img_A, img_B):
        x = torch.cat((img_A, img_B), 1)
        return self.model(x)

# =====
# BASELINE CNN (NO GAN)
# =====
class BaselineCNN(nn.Module):
    def __init__(self):
        super().__init__()
        self.encoder = nn.Sequential(
            nn.Conv2d(1, 64, 4, 2, 1),
            nn.ReLU(),
            nn.Conv2d(64, 128, 4, 2, 1),
            nn.ReLU()
        )
        self.decoder = nn.Sequential(
            nn.ConvTranspose2d(128, 64, 4, 2, 1),
            nn.ReLU(),
            nn.ConvTranspose2d(64, 1, 4, 2, 1),
            nn.Tanh()
        )
    def forward(self, x):
        return self.decoder(self.encoder(x))

# =====
# INITIALIZE
# =====
generator = UNetGenerator().to(DEVICE)
discriminator = PatchGANDiscriminator().to(DEVICE)
baseline = BaselineCNN().to(DEVICE)

```

```

=====
optimizer_G = optim.Adam(generator.parameters(), lr=Lr, betas=(0.5, 0.999))
optimizer_D = optim.Adam(discriminator.parameters(), lr=Lr, betas=(0.5, 0.999))
optimizer_base = optim.Adam(baseline.parameters(), lr=Lr)

adversarial_loss = nn.BCEWithLogitsLoss()
l1_loss = nn.L1Loss()

# =====
# TRAINING
# =====
print("Training Started...")

for epoch in range(EPOCHS):
    for edge_img, real_img in train_loader:

        edge_img = edge_img.to(DEVICE)
        real_img = real_img.to(DEVICE)

        valid = torch.ones((edge_img.size(0), 1, 7, 7), device=DEVICE)
        fake = torch.zeros((edge_img.size(0), 1, 7, 7), device=DEVICE)

        # ---- Train Generator ----
        optimizer_G.zero_grad()
        fake_img = generator(edge_img)
        pred_fake = discriminator(edge_img, fake_img)
        loss_GAN = adversarial_loss(pred_fake, valid)
        loss_pixel = l1_loss(fake_img, real_img)
        loss_G = loss_GAN + LAMBDA_L1 * loss_pixel
        loss_G.backward()
        optimizer_G.step()

        # ---- Train Discriminator ----
        optimizer_D.zero_grad()
        pred_real = discriminator(edge_img, real_img)
        loss_real = adversarial_loss(pred_real, valid)

        pred_fake = discriminator(edge_img, fake_img.detach())
        loss_fake = adversarial_loss(pred_fake, fake)

        loss_D = 0.5 * (loss_real + loss_fake)
        loss_D.backward()
        optimizer_D.step()

        # ---- Train Baseline ----
        optimizer_base.zero_grad()
        base_out = baseline(edge_img)
        loss_base = l1_loss(base_out, real_img)
        loss_base.backward()
        optimizer_base.step()

    print(f"Epoch {(epoch+1)}/{EPOCHS} D:{loss_D.item():.4f} G:{loss_G.item():.4f} Base:{loss_base.item():.4f}")

print("Training Completed")

```



